


# Solving problems by searching



Sanja Lazarova-Molnar, PhD

# Search as Problem-Solving Strategy

- many problems: reach goal state from a given starting point

# Examples

- getting from home to SDU
  - start: home
  - goal: SDU
  - operators: move one block, turn
- loading a moving truck
  - start: apartment full of boxes and furniture
  - goal: empty apartment, all boxes and furniture in the truck
  - operators: select item, carry item from apartment to truck, load item
- getting settled
  - start: items randomly distributed over the place
  - goal: satisfactory arrangement of items
  - operators: select item, move item

# Motivation


- search strategies - important methods for many approaches to problem-solving
- to use search - abstract formulation of the problem and the available steps
- search algorithms - basis for many optimization and planning methods

# Objectives

- formulate appropriate problems as search tasks
  - states, initial state, goal state, successor functions (operators), cost
- know the fundamental search strategies and algorithms
  - uninformed search
    - breadth-first, depth-first, uniform-cost, iterative deepening
  - informed search
    - best-first (greedy,  $A^*$ ), heuristics
- evaluate the suitability of a search strategy for a problem
  - completeness, optimality

# Search

- We will consider the problem of designing **goal-based agents** in **observable, deterministic, discrete, known** environments
- Example:

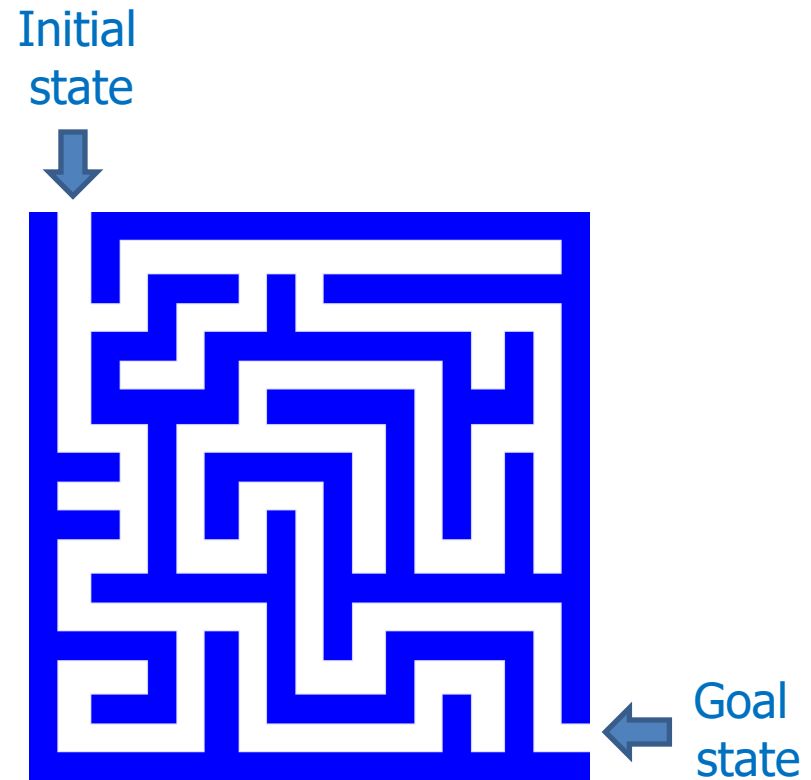


# Search

- We will consider the problem of designing **goal-based agents** in **observable, deterministic, discrete, known** environments
  - Solution: fixed sequence of actions
  - Search: process of looking for the sequence of actions that reaches the goal
  - Agent can ignore percepts during execution

# Search problem components

- **Initial state**
- **Successor Function (Operator)**
  - Result of doing an action in a state
- **Goal state**
- **Path cost**
  - Assume that it is a sum of nonnegative *step costs*



- **Optimal solution:** sequence of actions with the **lowest path cost** for reaching the goal



# Example: Romania

- On vacation in Romania; currently in Arad
- Flight leaves tomorrow from Bucharest

- **Initial state**

- E.g. Arad

- **Successor Function**

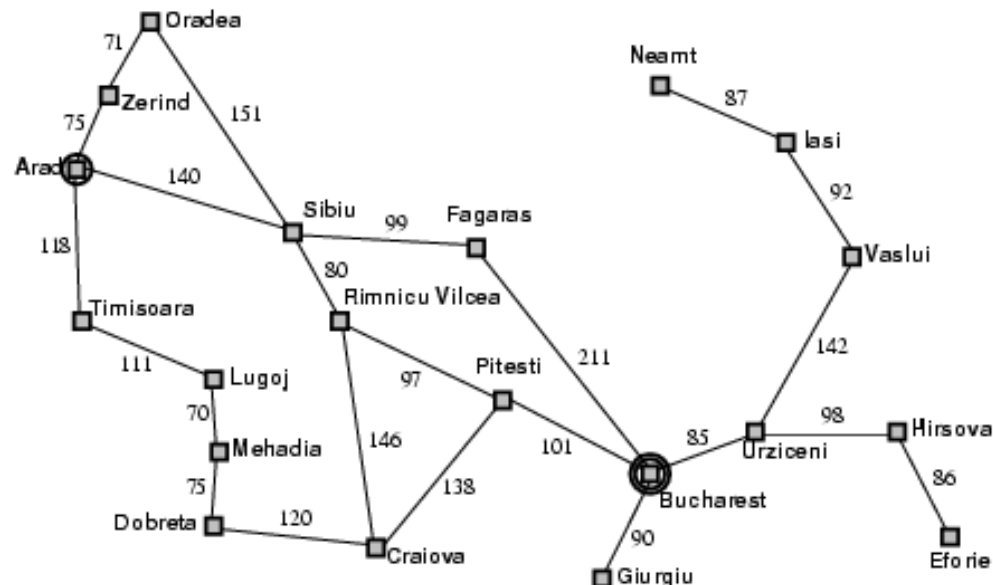
- $S(\text{Arad}) = \{\text{Zerind}, \text{Timisoara}, \text{Sibiu}\}$

- **Goal state**

- Bucharest

- **Path Cost**

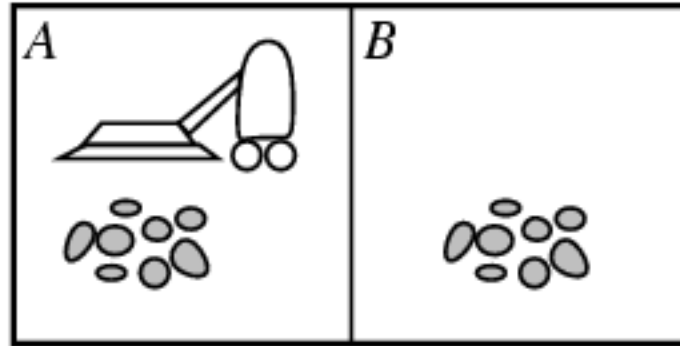
- Sum of edge costs



# State Space


- The initial state and the successor function define the **state space** of the problem
  - The set of all states reachable from initial state by any sequence of actions
  - Can be represented as a **directed graph** (nodes are states and links between nodes are actions)
- What is the state space for the Romania problem?

# Example: Vacuum world



- **Initial State**
  - Any
- **Goal State**
  - All clean
- **Successor function**
  - Described by state space (next slide)
- **Path Cost**
  - Could be the sum of the amounts of electricity consumed with each move

# Vacuum world state space graph



# Example: The 8-puzzle

- **Initial State**

- Any locations of tiles, number of states:

- 8-puzzle: 181,440 states
- 15-puzzle: 1.3 trillion states
- 24-puzzle:  $10^{25}$  states

- **Successor Function**

- Actions: Move blank left, right, up, down, and consequent states

- **Path cost**

- 1 per move

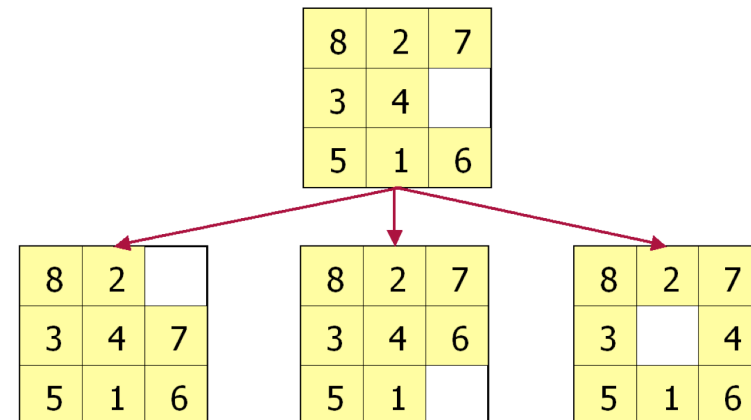
- Optimal solution of n-Puzzle is NP-hard

7	2	4
5		6
8	3	1



Start State

	1	2
3	4	5
6	7	8

Goal State




# Simpler: 3-Puzzle



# Building goal-based agents


## We must answer the following questions

- How do we represent the **state** of the “world”?
- What is the **goal** and how can we recognize it?
- What are the possible **actions**?
- What *relevant* information do we encode to describe states, actions and their effects and thereby solve the problem?



# Example: Route Planning

Find a route from Arad to Bucharest






# Example: The 8-Queens Puzzle

Place eight queens  
on a chessboard  
such that no queen  
attacks any other

We can generalize  
the problem to a  
 $N \times N$  chessboard



*What are the states, successor function, goal test?*

# Example: Remove 5 Sticks

## Remove 5 Sticks Problem


Remove exactly 5 of the 17 sticks so that the result forms exactly 3 squares



How could this be stated as a search problem? (States? Initial State? Successor Function? Goal Test? Path Cost?)

# Tree Search

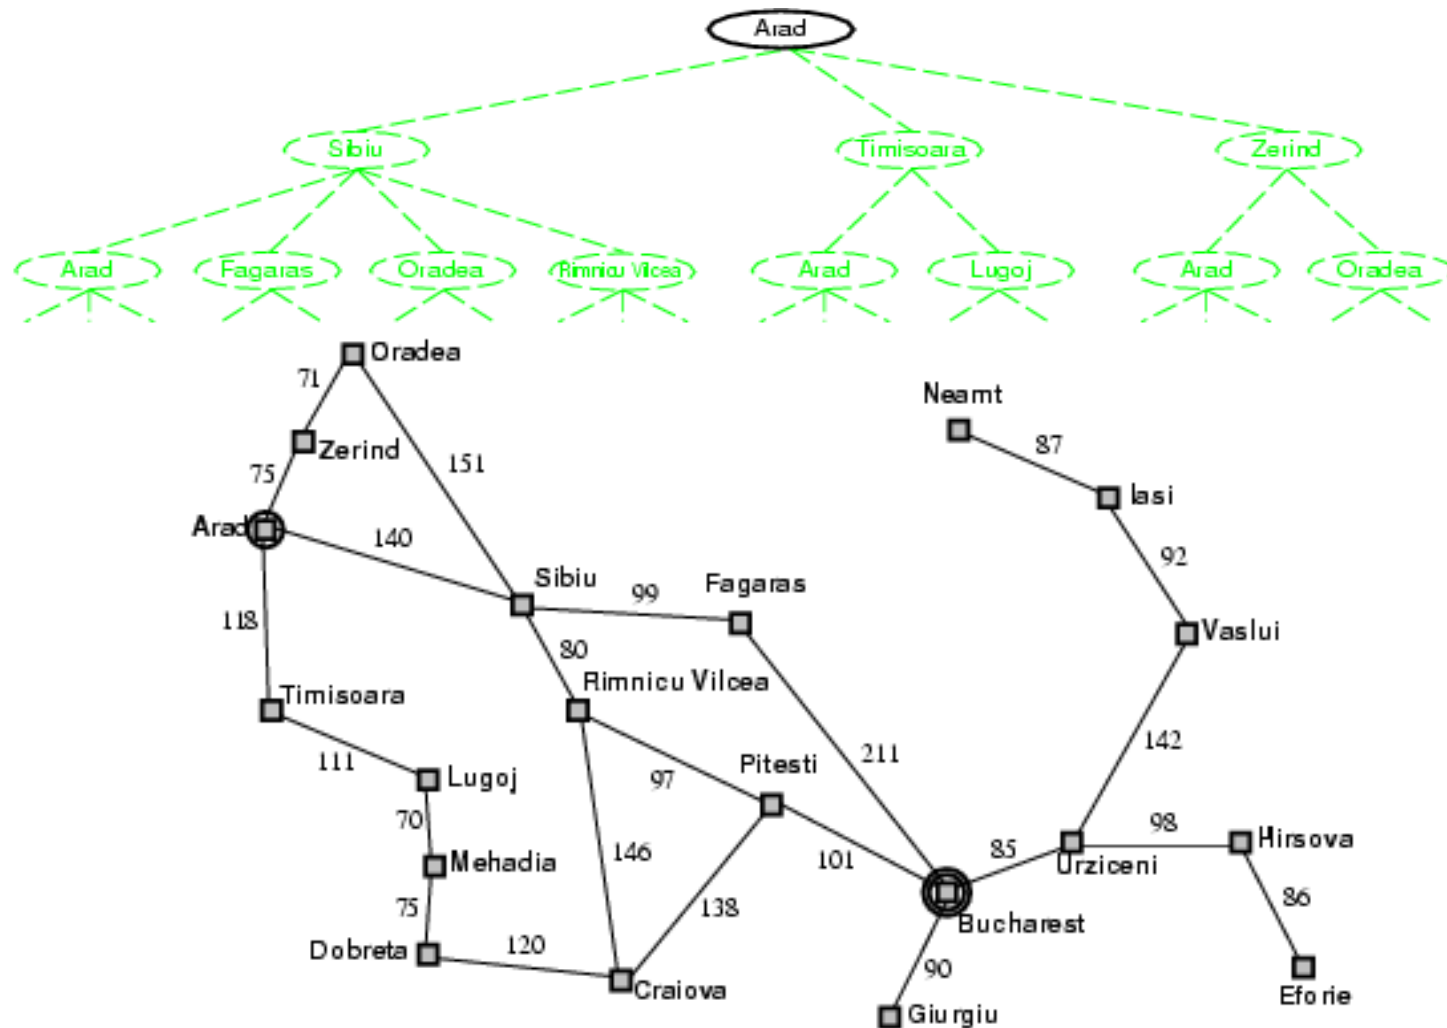
- Begin at the root node (starting state) and **expand** it by making a list of all possible successor states
- Maintain a **fringe** or a list of unexpanded states
- At each step, pick a state from the fringe to expand
- Keep going until you reach the goal state
- Try to expand as few states as possible
- A solution is a path ending in the goal state
- Nodes vs. states
  - A state is a representation of a physical configuration, while a node is a data structure that is part of the search tree



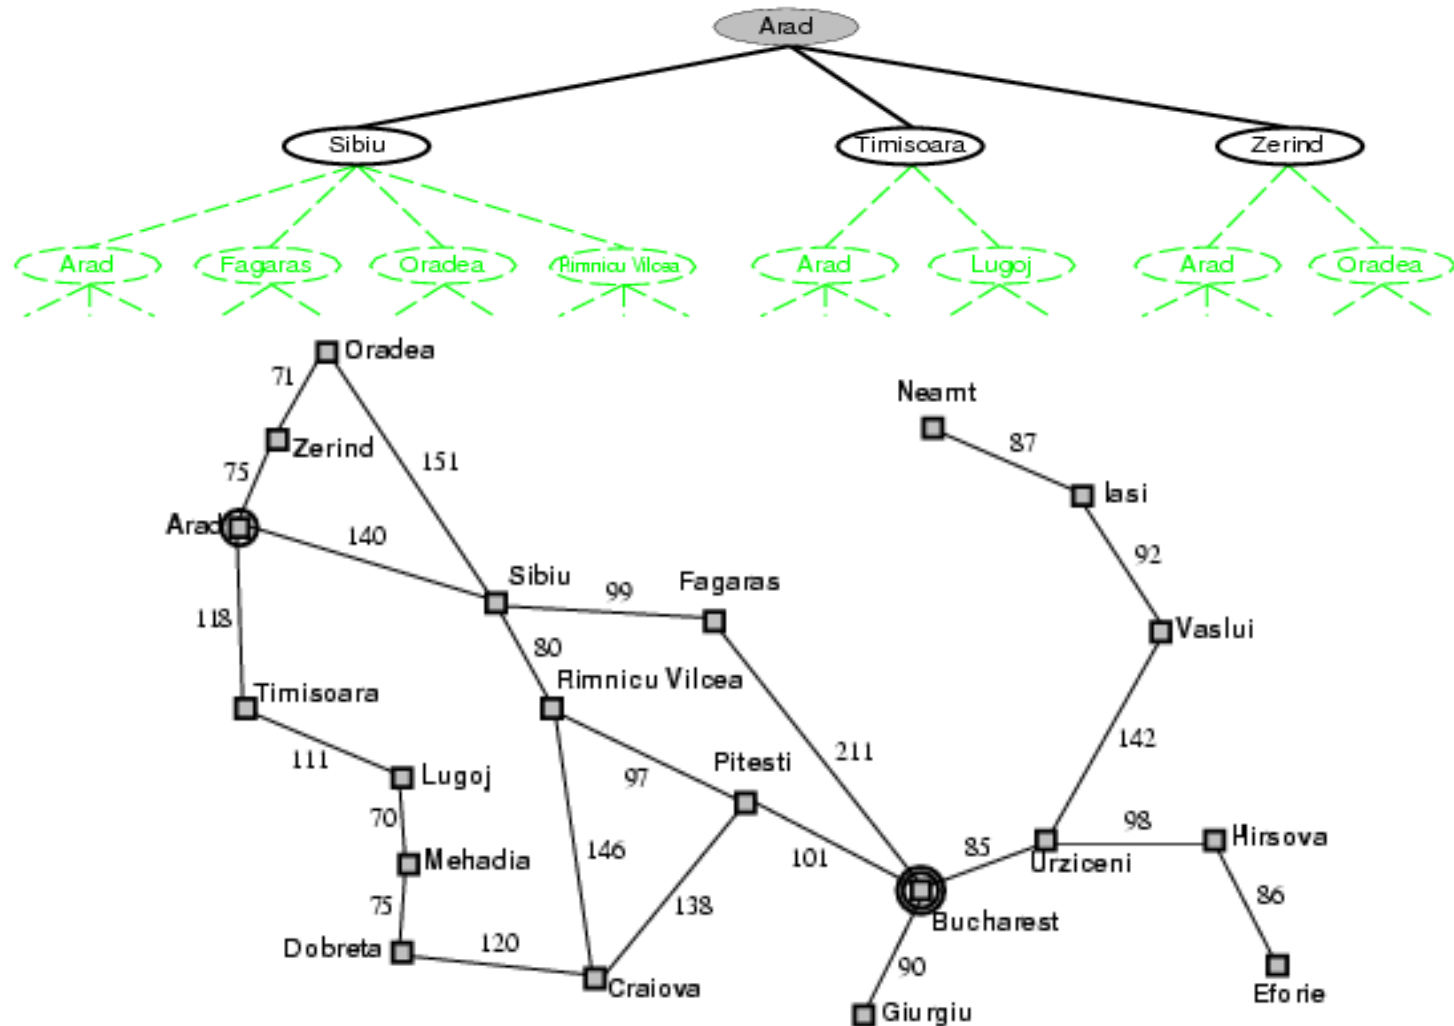
# Tree Search Algorithm Outline

- Initialize the **fringe** using the **starting state**
- While the fringe is not empty
  - Choose a fringe node to expand according to **search strategy**
  - If the node contains the **goal state**, return solution
  - Else **expand** the node and add its children to the fringe
- To handle repeated states:
  - Keep an **explored set**; add each node to the explored set every time you expand it
  - Every time you add a node to the fringe, check whether it already exists in the fringe with a higher path cost, and if yes, replace that node with the new one

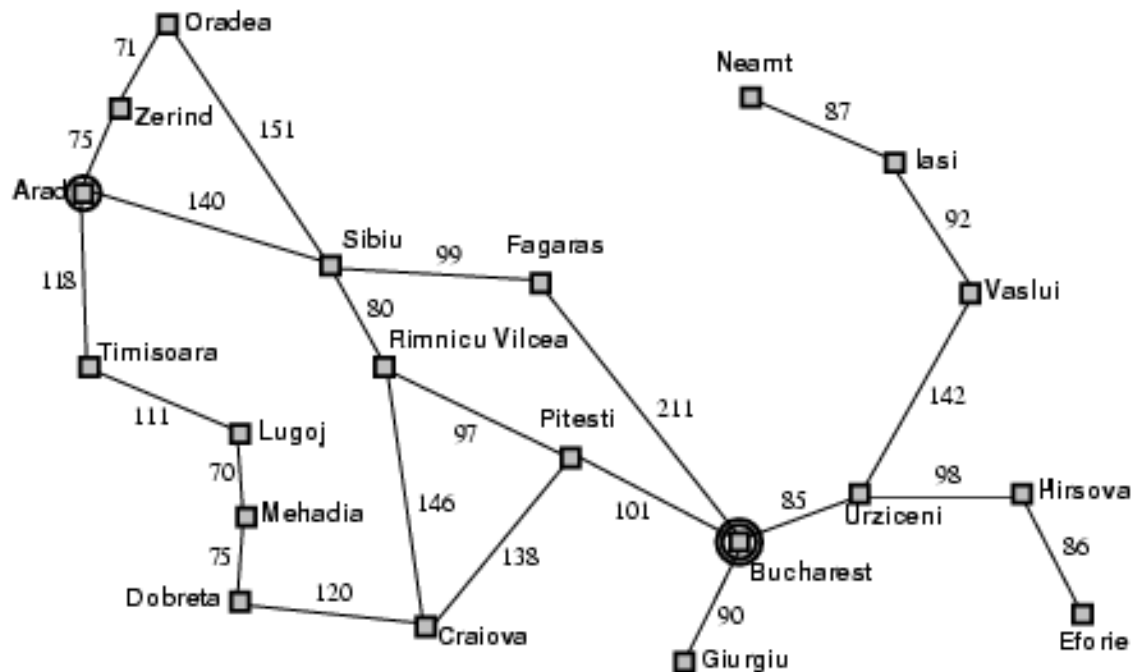

# Tree search example



# Tree search example



# Tree search example



# Search strategies

- A **search strategy** is defined by picking the order of node expansion
- Strategies are evaluated along the following dimensions:
  - **Completeness**: does it always find a solution if one exists?
  - **Optimality**: does it always find a least-cost solution?
  - **Time complexity**: number of nodes generated
  - **Space complexity**: maximum number of nodes in memory
- Time and space complexity are measured in terms of
  - $b$ : maximum branching factor of the search tree
  - $d$ : depth of the least-cost (cheapest) solution
  - $m$ : maximum length of any path in the state space (may be infinite)




# Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Iterative deepening search


# Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Properties of breadth-first search

- **Complete?**

Yes

- **Optimal?**

Yes – if cost = 1 per step

- **Time?**

Number of nodes in a  $b$ -ary tree of depth  $d$ :  $O(b^d)$   
( $d$  is the depth of the optimal solution)

- **Space?**

$O(b^d)$


- Space is the bigger problem (more than time)

# Uniform-cost search

- Expand least-cost unexpanded node
- Implementation: *fringe* ordered by path cost (priority queue)
- Equivalent to breadth-first if step costs all equal
- **Complete?**  
Yes
- **Optimal?**  
Yes – nodes expanded in increasing order of path cost
- **Time?**  
Number of nodes with path cost  $\leq$  cost of optimal solution ( $C^*$ ),  $O(b^{C^*/\epsilon})$ ,  
where every step cost is at least  $\epsilon > 0$   
This can be greater than  $O(b^d)$ : the search can explore long paths consisting  
of small steps before exploring shorter paths consisting of larger steps
- **Space?**  
 $O(b^{C^*/\epsilon})$



# Uniform-Cost Search Example



Start Node: A

Goal Node: G

*Step 1*

Fringe:

**Node A**

**Cost 0**

Explored: -

*Step 2*

Expand A

Fringe:

**Node D B**

**Cost 3 5**

Explored: A

*Step 3*

Expand D

Fringe:

**Node B E F**

**Cost 5 5 5**

Explored: A D

*Step 4*

Expand B

Fringe:

**Node E F C**

**Cost 5 5 6**

Explored: A D B

*Step 5*

Expand E

Fringe:

**Node F C**

**Cost 5 6**

Explored: A D B E

*Step 6*

Expand F

Fringe:

**Node C G**

**Cost 6 8**

Explored: A D B E F

*Step 7*

Expand C

Fringe:

**Node G**

**Cost 8**

Explored: A D B E F C

*Step 8*


Expand G

Found the path:

A to D to F to G


# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search


- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front






# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

- Expand deepest unexpanded node
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Properties of depth-first search

- **Complete?**

Fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

→ complete in finite spaces

- **Optimal?**

No – returns the first solution it finds

- **Time?**

Could be the time to reach a solution at maximum depth  $m$ :  $O(b^m)$

Terrible if  $m$  is much larger than  $d$

But if there are lots of solutions, may be much faster than BFS

- **Space?**

$O(bm)$ , i.e. linear space!

# Iterative deepening search

- Use DFS as a subroutine
  1. Check the root
  2. Do a DFS searching for a path of length 1
  3. If there is no path of length 1, do a DFS searching for a path of length 2
  4. If there is no path of length 2, do a DFS searching for a path of length 3...



# Iterative deepening search

Limit = 0




# Iterative deepening search

Limit = 1




# Iterative deepening search

Limit = 2



# Iterative deepening search

Limit = 3





# Properties of iterative deepening search

- **Complete?**

Yes

- **Optimal?**

Yes, if step cost = 1

- **Time?**

$$(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$$

- **Space?**

$$O(bd)$$

# Next Class

- Informed Search